CS2030 Lecture 12

From Sequential to Parallel Streams

Henry Chia (hchia@comp.nus.edu.sg)

Semester 2 2019 / 2020

Lecture Outline

☐ Concurrency versus parallelism

☐ Parallel streams

  – Debugging parallel streams

  – Comparing sequential and parallel streams

☐ Correctness of parallel streams

  – `reduce` operator

  – Accumulator and combiner

☐ Overhead of parallelization

Concurrency vs Parallelism

- A single core processor executes one instruction at a time

    - Only one process can run at any one time
    - Context-switching allows multi-tasking on a single processor

- Concurrent programs run concurrently via threads

    - OS switches between threads
    - Separate unrelated tasks into separate threads
    - Improves processor utilization

- Parallel computing involves multiple subtasks running at the same time on multiple (possibly multi-core) processors
- Parallel programs are concurrent, but not all concurrent programs are parallel

Sequential

Concurrent

thread

Parallel

Parallel Streams

- ☐ Parallel streams use a common `ForkJoinPool` via the static `ForkJoinPool.commonPool()` method

  ```
  jshell> ForkJoinPool.commonPool().getParallelism()
  $1 ==> 23
  ```

- ☐ The level of parallelism can be controlled by setting the system property during program run

  ```
  System.setProperty("java.util.concurrent.ForkJoinPool.common.parallelism", "4")
  ```

  or including the following flag when running the program

  ```
  $ java -Djava.util.concurrent.ForkJoinPool.common.parallelism=4 ...
  ```

- ☐ Similar to the method `stream()`, Java `Collection`(s) also support the method `parallelStream()` to create a parallel stream of elements

Parallel Streams

□ Using prime number testing as an example

```java
static boolean isPrime(int n) {
    return IntStream
        .rangeClosed(2, (int) Math.sqrt(n))
        .noneMatch(x -> n % x == 0);
}
```

□ Count number of primes between $2,000,000$ and $3,000,000$

```java
long count = IntStream.range(2_000_000, 3_000_000)
    .filter(x -> isPrime(x))
    .count();
```

stream

is prime?

Parallel Streams

☐ Parallelizing the seach for primes

```
long count = IntStream.range(2_000_000, 3_000_000)
    .parallel()
    .filter(x -> isPrime(x))
    .count();
```

☐ The `parallel()` intermediate operation turns on a boolean flag that switches the stream pipeline to be parallel

– Invoked anywhere between the data source and terminal
– The counter operation is `sequential()`

```java
public static void main(String[] args) {
    if (args.length != 0) {
        System.setProperty(
            "java.util.concurrent.ForkJoinPool.common.parallelism",
            args[0]);
    }
    System.out.println("Number of worker threads: " +
            ForkJoinPool.commonPool().getParallelism());

    Instant start = Instant.now();
    long howMany = IntStream.range(2_000_000, 3_000_000)
        .parallel()
        .filter(x -> isPrime(x))
        .count();
    Instant stop = Instant.now();

    System.out.println(howMany + " : " +
            Duration.between(start, stop).toMillis() + "ms");
}
```

Debugging Parallel Streams

☐ To time the execution of a process,

– `java.time.Instant`'s `now()` method returns the current `Instant` from the system clock

– `java.time.Duration`'s `between()` returns the `Duration` of two `Instances` (an implementation of `Temporal`)

– `Duration`'s `toMillis()`/`toNanos()`/... extracts the desired representation of the duration

```java
java.util.Instant;
java.util.Duration;

Instant start, stop;
start = Instant.now();
/* perform some task */
stop = Instant.now();

long timeInMillis = Duration.between(start, stop).toMillis();
```

Debugging Parallel Streams

☐ To debug and manage each execution thread

- <mark>Thread.currentThread()</mark> (or <mark>Thread.currentThread().getName()</mark>) to retrieve the identity of the thread

- Thread.sleep(**long** millis) causes the currently executing thread to sleep (i.e. temporarily cease execution) for the specified number of milliseconds

  ▷ Used within a **try.. catch** block
  ▷ Example, letting a thread sleep for one second

```
try {
    ...
    Thread.sleep(1000);
    ...
} catch (InterruptedException e) { }
```

Debugging Parallel Streams

□ Effect of parallelizing a stream

```java
int sum = IntStream.of(1, 2, 3, 4, 5)
    .parallel()
    .filter(x -> {
        System.out.println("filter:  " + x + " "
            + Thread.currentThread().getName());
        return x % 2 == 1;
    })
    .map(x -> {
        System.out.println("map:     " + x + " "
            + Thread.currentThread().getName());
        return x;
    })
    .reduce(0, (x, y) -> {
        System.out.println("reduce:  " + x + " + " + y + " "
            + Thread.currentThread().getName());
        return x + y;
    });
System.out.println(sum);
```

Comparing Sequential and Parallel Streams

☐ Suppose given the following task unit

```java
class OneSecondTask {
    int ID;

    public OneSecondTask(int ID) {
        this.ID = ID;
    }

    public int compute() {
        System.out.println(Thread.currentThread().getName());
        try {
            Thread.sleep(1000);
        } catch (InterruptedException e) {
            throw new RuntimeException(e);
        }
        return ID;
    }
}
```

☐ Comparing sequential and parallel streams:

```java
Stream<OneSecondTask> tasks = IntStream.range(0, 10)
        .mapToObj(x -> new OneSecondTask(x));
runTasks(tasks);
```

Comparing Sequential and Parallel Streams

```java
public static void runTasks(Stream<OneSecondTask> tasks) {
    Instant start = Instant.now();
    List<Integer> result = tasks
        .map(x -> x.compute())
        .collect(Collectors.toList());
    Instant stop = Instant.now();
    System.out.print(result + " ");
    System.out.println(Duration.between(start,stop).toMillis() + "ms");
}
```

☐ Sequential stream on 4 worker threads:

```
main
main
main
main
main
main
main
main
main
main
[0, 1, 2, 3, 4, 5, 6, 7, 8, 9] 10003ms
```

Comparing Sequential and Parallel Streams

```
Stream<OneSecondTask> tasks = IntStream.range(0, 10)
    .mapToObj(x -> new OneSecondTask(x));
runTasks(tasks.parallel());
```

☐ Parallel stream on 4 worker threads:

```
main
ForkJoinPool.commonPool-worker-1
ForkJoinPool.commonPool-worker-3
ForkJoinPool.commonPool-worker-2
ForkJoinPool.commonPool-worker-4
ForkJoinPool.commonPool-worker-3
ForkJoinPool.commonPool-worker-2
ForkJoinPool.commonPool-worker-4
ForkJoinPool.commonPool-worker-1
ForkJoinPool.commonPool-worker-3
[0, 1, 2, 3, 4, 5, 6, 7, 8, 9] 3006ms
```

Correctness of Parallel Streams

□ To ensure correct parallel execution, stream operations

– must not interfere with stream data (true for sequential streams also)

– preferably stateless with no side effects

□ Example of interference:

```java
List<String> list = new ArrayList<>(
        List.of("abc", "def", "xyz"));

list.stream()
    .peek(str -> {
        if (str.equals("xyz")) {
            list.add("pqr");
        }
    })
    .forEach(x -> {});
```

Correctness of Parallel Streams

☐ Example of side-effects:

```
List<Integer> list = new ArrayList<>(
        Arrays.asList(1, 3, 5, 7, 9, 11, 13, 15, 17, 19));
List<Integer> result = new ArrayList<>();
```

☐ The following is erroneous

```
list.parallelStream()
    .filter(x -> isPrime(x))
    .forEach(x -> result.add(x));
```

☐ Use `.collect` instead

```
result = list.parallelStream()
        .filter(x -> isPrime(x))
        .collect(Collectors.toList());
```

☐ Can also consider using `forEachOrdered()`, or a thread-safe list `CopyOnWriteArrayList`

forEachOrder ensure that all elements are there, but not in order

Inherently Parallelizable **reduce**

□ Consider `Stream`'s three-argument `reduce` method:

```
<U> U reduce(U identity,
        BiFunction<U,? super T,U> accumulator,
        BinaryOperator<U> combiner)
```

□ Rules to follow when parallelizing

– `combiner.apply(identity, i)` must be equal to `i`
– `combiner` and `accumulator` must be associative, i.e. order of application does not matter
– `combiner` and `accumulator` must be compatible, i.e. `combiner.apply(u, accumulator.apply(identity, t))` must be equal to `accumulator.apply(u, t)`

Accumulator and Combiner

☐ The following example complies with the above rules:

```
Stream.of(1,2,3,4)
    .parallel()
    .reduce(1, (x,y) -> x * y, (x,y) -> x * y)
```

☐ To see the effects of accumulator and combiner:

```
sum = Stream.of(1, 2, 3, 4, 5)
    .parallel()
    .filter(x -> {
        System.out.println("filter:  " + x + " "
            + Thread.currentThread().getName());
        return x % 2 == 1;
    })
    .reduce(0,
        (x, y) -> {
            System.out.println("accumulate: " + x + " + " + y + " " +
                    Thread.currentThread().getName());
            return x + y;
        },
        (x, y) -> {
            System.out.println("combine: " + x + " + " + y + " "
                    + Thread.currentThread().getName());
            return x + y;
        }
    );
```

Accumulator and Combiner

□ High-level schematic on parallelizing a stream

accumulate

combine

Accumulator and Combiner

□ Output from a sample run:

```
filter:  5 ForkJoinPool.commonPool-worker-1
filter:  4 ForkJoinPool.commonPool-worker-3
filter:  1 ForkJoinPool.commonPool-worker-3
filter:  3 main
filter:  2 ForkJoinPool.commonPool-worker-2
accumulate: 0 + 5 ForkJoinPool.commonPool-worker-1 // (A)
accumulate: 0 + 1 ForkJoinPool.commonPool-worker-3
combine: 1 + 0 ForkJoinPool.commonPool-worker-3    // (C)
accumulate: 0 + 3 main
combine: 0 + 5 ForkJoinPool.commonPool-worker-1
combine: 3 + 5 ForkJoinPool.commonPool-worker-1    // (B)
combine: 1 + 8 ForkJoinPool.commonPool-worker-1
9
```

(A)   Accumulation with identity

(B)   Combining results from accumulation

(C)   Combining with identity

Accumulator and Combiner

☐ Erroneous examples where rules are violated

- – `combiner.apply(identity, i)` not equal to `i`

```
double result = Stream
    .of(1, 2, 3, 4)
    .parallel()
    .reduce(1.0,
            (x, y) -> x + y,
            (x, y) -> x + y);
```

- – for division, the order of application **does** matter

```
double result = Stream
    .of(1, 2, 3, 4)
    .parallel()
    .reduce(24.0,
            (x, y) -> 1.0 * x / y,
            (x, y) -> 1.0 * x / y);
```

Fork and Join in Parallel Streams

- □ <mark>`parallel()` runs `fork` to create sub-tasks running the same chain of operations on sub-streams</mark>

    - Processes for sub-tasks are run in multiple threads when appropriate
    - Threads are shared from a common **Fork Join Pool**

- □ `combiner` in `reduce` runs `join` to combine the results

Overhead in Parallel Streams

- Should we exploit parallelism to the fullest?

```java
return IntStream
    .rangeClosed(2, (int) Math.sqrt(n))
    .parallel()
    .noneMatch(x -> n % x == 0);
```

- Parallelizing a trivial task actually creates more work in terms of parallelizing overhead
- Parallelization is worthwhile only if the task is complex enough that the benefit of parallelization outweights the overhead

  - In primality testing, checking (n % x == 0) is trivial;
  - Parallelizing it induces more overhead in terms of processing the forks and joins

- Holds true for all parallel and concurrent programs

Lecture Summary

☐ Familiarity with the use of sequential and parallel streams

☐ Able to compare performances between sequential and parallel streams

☐ Able to debug parallel streams

☐ Adherence to rules for parallelizing streams

☐ Appreciate the overhead involved in parallelizing streams